DTU Library

Two Methods for Antialiased Wireframe Drawing with Hidden Line Removal

Bærentzen, Jakob Andreas; Munk-Lund, Steen; Gjøl, Mikkel; Larsen, Bent Dalgaard

Published in:
Proceedings of the Spring Conference in Computer Graphics

Publication date:
2008

Link back to DTU Orbit

Citation (APA):
Bærentzen, J. A., Munk-Lund, S., Gjøl, M., & Larsen, B. D. (2008). Two Methods for Antialiased Wireframe Drawing with Hidden Line Removal. In *Proceedings of the Spring Conference in Computer Graphics*

Two Methods for Antialiased Wireframe Drawing with Hidden Line Removal

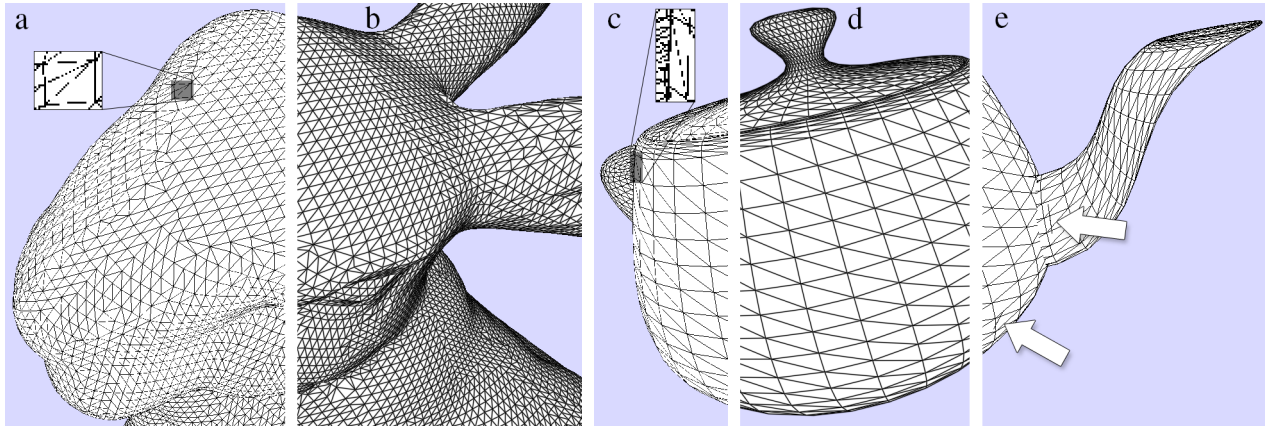J. Andreas Bærentzen, Steen Lund Nielsen, Mikkel Gjøl, and Bent D. Larsen*

Figure 1: The left image shows Stanford Bunny drawn using (a) the offset method with a constant offset, and (b) our single pass method. The right image shows the Utah Teapot drawn using (c) a constant offset, (d) our ID buffer method and (e) a slope dependent offset. The white arrows indicate artifacts introduced by the slope dependent offset. Notice the stippling in the constant offset images.

Abstract

Two novel and robust techniques for wireframe drawing are proposed. Neither suffer from the well-known artifacts associated with the standard two pass, offset based techniques for wireframe drawing. Both methods draw prefiltered lines and produce high-quality antialiased results without super-sampling. The first method is a single pass technique well suited for convex N-gons for small N (in particular quadrilaterals or triangles). It is demonstrated that this method is more efficient than the standard techniques and ideally suited for implementation using geometry shaders. The second method is completely general and suited for arbitrary N-gons which need not be convex. Lastly, it is described how our methods can easily be extended to support various line styles.

1 Introduction

Modern graphics cards are machines for drawing textured triangles, and even though they have become extremely sophisticated and completely programmable, they are very poor at drawing lines. While line drawing is not so often required in "consumer" graphics, wireframe drawing is commonly used by graphics professionals (artists and programmers alike) to gauge the quality of polygonal meshes with regard to various uses, and removal of the hidden lines makes it much easier to understand the geometry. Therefore, it is unsurprising that there are well known, standard procedures for wireframe drawing with hidden line removal [McReynolds and Blythe 1999]. However, although almost universally used, these methods, which were designed for fixed function graphics

pipelines, suffer from inherent problems with two types of artifacts, namely stippling where a visible line is intermittently invisible (cf. Figure 1a,c) and disocclusion of lines that should be hidden (cf. Figure 1e). Furthermore, these methods are based on the hardware supported line primitive which tends to be costly to render, especially because the quality is very low unless hardware antialiasing is enabled.

In this paper, we propose two new algorithms, which are both designed for modern, commodity graphics hardware supporting fragment and vertex shaders. The algorithms are designed with the goals that they should

- support the combination of wireframe and filled polygons.

- not suffer from the artifacts of the standard techniques.

- produce high quality, antialiased lines (without hardware super sampling).

Our first solution is a simple, single pass method which solves the hidden line removal problem by drawing the wireframe mesh as a part of the rendering of the filled polygons. At little extra effort, the wireframe is drawn using smooth, prefiltered lines [Gupta and Sproull 1981]. The single pass method is shown in Figure 1b and discussed in Section 2. For hardware supporting *geometry shaders*, a very efficient implementation is possible.

The single pass method is most efficient for convex N-gons for relatively small N. For completely general polygons (not necessarily planar, convex and of arbitrary N) we propose a second method, which we denote the identity (ID) buffer method. This method, which also produces artifact free, prefiltered lines, is shown in Figure 1d and discussed in Section 3. Both our methods can be extended to handle a variety of line styles. Most importantly, we can distance attenuate the line thickness, but many other possibilities present themselves as discussed in Section 4 where our results are presented. Below, we discuss related work, and in the final section (Section 5) we draw conclusions and point to future work.

---

*Andreas Bærentzen is with the Informatics and Mathematical Modelling Department of the Technical University of Denmark. Email: jab@imm.dtu.dk. Steen Lund Nielsen is with Deadline Games. Email: steen@munk-lund.dk. Mikkel Gjøl is with Zero Point Software. Email: pixelmager@gmail.com. Bent D. Larsen is with Dalux and the Department of Informatics and Mathematical Modelling, Technical University of Denmark. Email: bdl@dalux.dk

1.1 Related Work

The standard wireframe rendering techniques are described well in [McReynolds and Blythe 1999]. There is a number of variations, but they all have in common that the filled polygons and the polygon edges are drawn in separate passes. Assuming that the filled polygons have been rendered to the depth buffer, the depth value of a fragment generated from a line (polygon edge) should be either identical to the stored depth value (from a filled polygon) or greater – depending on whether the line is visible or hidden. In reality it is less simple, because the line depth values depend only on the end points whereas the depth value of a fragment that originates from a filled polygon depends on the slope of the polygon. Thus, a fragment produced by polygon rasterization can have a very different depth value from a fragment at the same screen position produced from the line rasterization of one of its edges.


Figure 2: Wireframe drawing using the offset methods (b,d) and our ID buffer method (c) which employs prefiltered lines (and no other method of antialiasing). In (a) no offset has been used, and stippling artifacts are caused by the steep slope of the gray triangle. In (b) a large constant offset (-30000) is used but the result still suffers from stippling artifacts although the edge of the gray triangle penetrates the blue triangle. In (d) a slope dependent offset is used which fixes the stippling but causes disocclusion.

Stippling artifacts where the lines are partly occluded are caused by this issue (Figure 2a). To solve this problem, a depth offset or a shifting of the depth range is used to bias the depth values from either the polygons or the lines to ensure that the lines are slightly in front of the polygons. However, this depth offset has to be very large to avoid stippling in the case of polygons with steep slopes. Hence, in most 3D modeling packages, stippling can still be observed along the edges of polygons with steep slopes. In other cases, the offset is too large and hidden parts of the wireframe become visible. In fact, as shown in Figure 2b, the offset can be both too large and too small for the same polygon. The OpenGL API provides a slope-dependent offset, but this often yields offsets which are too great leading to (usually sporadic) disocclusion of hidden lines as seen in Figure 2d. In the following, we will refer to these methods as *offset methods*.

A few authors have tried to address the issues in wireframe drawing. Herrell et al. [Herrell et al. 1995] proposed the edging-plane technique. Essentially, this technique ensures that a fragment from a polygon edge is drawn if the corresponding fragment of the filled polygon passed the depth test. This method can be implemented either using special hardware or using the stencil buffer. In any

case frequent state changes are required for this method. A variant of the method is discussed in [McReynolds and Blythe 1999]. In Section 3 we propose a similar but more efficient algorithm that requires few state changes. Wang et al. [Wang et al. 1999] proposed a different technique which explicitly computes occlusion relationships between polygons. While elegant this method cannot be implemented on a GPU and there is no simple way to handle intersecting polygons. Finally, in a method which is most practical for pure quad meshes, Rose et al. [Rose and Ertl 2003] encoded signed line distances in textures in order to be able to texture map a detailed mesh onto a simplified model.

In computer graphics, antialiasing is normally performed by taking multiple samples per pixel. These samples are subsequently combined using a simple discrete filter e.g. a straight average [Akenine-Möller and Haines 2002]. This method is always feasible to implement, but often many samples are required to obtain reasonable quality. Far better quality is obtained if continuous filtering is used. However, this is only possible if we can directly evaluate the convolution of a continuous image with a low pass filter at the point where we desire to compute the intensity of the filtered image. Clearly, this is only possible for filters of finite support, but in the case of lines, Gupta and Sproull [Gupta and Sproull 1981] observed that the convolution depends only on the distance to the line. Thus, the convolution can be implemented using a 1D lookup table. McNamara et al. [McNamara et al. 2000] proposed an implementation of prefiltered line drawing involving small changes to the fragment generation part of a graphics card, and, more recently, Chan et al. [Chan and Durand 2005] have proposed a method which is implementable on modern graphics cards using fragment shaders. For our second wireframe method (cf. Section 3), we use a similar technique for prefiltered line drawing.

This work is related to *Non–Photorealistic Rendering* (NPR) in two ways: First of all, both our methods allow for variations of the line style which is also a common feature in NPR systems. Secondly, line drawing (especially silhouette and crease lines) is a central problem in NPR rendering, and wireframe rendering and NPR share the problem of determining the visibility of lines. See [Isenberg et al. 2003] for a survey of line drawing in NPR. We use a variation of the ID buffer method from Markosian [Markosian 2000] in order to resolve visibility in our second wireframe method.

The first of the two methods described here was originally presented in a SIGGRAPH sketch [Bærentzen et al. 2006] based on which NVIDIA implemented a Direct3D 10 version of the method for their SDK [Gateau 2007]. However, this is the first reviewed publication on the first method, and nothing has been reported on the second method.

2 The Single Pass Method

In the single pass method, the polygon edges are drawn as an integral part of drawing the filled polygons. For each fragment of a rasterized polygon, we need to compute the line intensity, $I(d)$, which is a function of the window space distance, d, to the boundary of the polygon (see Figure 3) which we assume is convex. The fragment color is computed as the linear combination $I(d)C_l + (1 - I(d))C_f$ where C_l and C_f are the line and face colors, respectively.

Note that a convex N-gon is simply the intersection of N half-planes and that d can be computed as the minimum of the distances to each of the N lines which delimit the half-planes. In 2D the distance to a line is an *affine* function. Now, suppose, we are given a triangle in 2D window space coordinates, $\mathbf{p}_i$, $i \in \{1, 2, 3\}$. Let $d_i = d(\mathbf{p}_i)$ be the distance from $\mathbf{p}_i$ to a given line in window space coordinates

$$I_{\mathbf{p}} = I(d)$$
$$= I(\min(d_1, d_2, d_3))$$

$I(d)$

Figure 3: The intensity, $I_{\mathbf{p}}$, of a fragment centered at $\mathbf{p}$ is computed based on the window space distances d_1, d_2, and d_3 to the three edges of the triangle.

and let α_i be the barycentric coordinates of a point $\mathbf{p}$ with respect to $\mathbf{p}_i$. d is an affine function of $\mathbf{p}$ and $\sum_i \alpha_i = 1$, it follows that

$$d(\mathbf{p}) = d(\alpha_1 \mathbf{p}_1 + \alpha_2 \mathbf{p}_2 + \alpha_3 \mathbf{p}_3) = \alpha_1 d_1 + \alpha_2 d_2 + \alpha_3 d_3 \quad ,$$

which has the important implication that we can compute the distance to the line per vertex and interpolate the distance to every fragment using linear interpolation in window space. In the case of triangles, we only need to compute the distance to the opposite edge for every vertex since it shares the two other edges with its neighboring vertices. For quadrilaterals we need the distance to two lines and so on. The distance from a vertex i to the line spanned by the segment jk is computed

$$d_i^{jk} = \frac{|(\mathbf{p}_j - \mathbf{p}_i) \times (\mathbf{p}_k - \mathbf{p}_i)|}{\|\mathbf{p}_k - \mathbf{p}_j\|} \tag{1}$$

2.1 Technical Details

There used to be a slight problem; graphics hardware performs perspective correct interpolation rather than linear interpolation. This problem is fixed by the latest graphics standards (Direct3D and OpenGL) which define a simple linear interpolation mode. However, even when this is not available in the API or supported by the hardware, there is a simple solution: For any datum that requires linear interpolation, we multiply the datum by the vertex w value before rasterization and multiply the interpolated datum by the interpolated $1/w$ in the fragment stage. This simple procedure negates the effects of perspective correction, and the trick works for arbitrary polygons: Although different modes of interpolation may be used [Segal and Akeley 2004], the interpolation is always linear (once perspective correction has been negated) thus affine functions must be reproduced.

In summary, for each vertex of an N-gon, before rasterization we need to

G1 compute the world to window space transformation of all vertices.

G2 compute the distance from vertex i to all edges.

G3 output an N vector $\mathbf{d}$ of distances for each vertex. For instance, for a triangle and a given vertex with index $i = 1$, this vector is $\mathbf{d}' = [0\, d_1^{23}\, 0]^T w_1$ where d_1^{23} is the distance from the first vertex to its opposite line. If non-perspective correct interpolation is not supported, each vector is premultiplied by the w value of vertex i.

After rasterization, i.e. in the fragment shader we receive the interpolated $\mathbf{d}$ and

R1 compute the minimum distance: $d = \min(d^{12}, d^{23}, ..., d^{n1})$ (having first multiplied $\mathbf{d}$ by the interpolated $1/w$ if required).

R2 output the fragment color: $C = C_l I(d) + (1 - I(d)) C_f$

Step **G1** is generally always performed in the vertex shader. Step **G2** should be performed in the geometry shader since the geometry shader has access to all vertices of the polygon (provided it is a triangle). Consequently, if the geometry shader is not used, the vertex shader needs to receive the other vertices as attributes in order to compute the distances. This entails that the application code is not invariant to whether the wireframe method is used, that additional data needs to be transmitted, and, finally, that the use of indexed primitives is precluded since the attributes of a vertex depend on which triangle is drawn. The appendix contains GLSL shader code for hardware which supports geometry shaders and disabling of texture correct interpolation.

One concern was raised in [Gateau 2007]. If one or two vertices of the polygon are behind the eye position in world coordinates, these vertices do not have a well-defined window space position. As a remedy, 2D window space representations of the visible lines are computed and sent to the fragment shader where the distances are computed. While this greatly complicates matters and makes the fragment shader more expensive, the case is rare which entails that the more complex code path is rarely taken.

2.2 Computing the Intensity

The convolution of a line and a symmetric filter kernel is a function of distance only in the case of infinite lines. Hence, other methods must be used to handle line segments and junctions of segments. To ensure that lines have rounded end points other authors combine several distance functions per line [McNamara et al. 2000], and to compose multiple lines, blending is sometimes used [Chan and Durand 2005]. In this work, we take a slightly different approach and compute the pixel intensity always as a a function, $I(d)$, of the distance, d, to the closest polygon edge.

Thus, our method does not perform prefiltering in the sense of exact continuous domain convolution of a line image and a kernel[1], and, in fact, I is not even a tabulation of convolutions. However, the choice of I is important; I should be a smooth, monotonically decreasing function such that $I(0) = 1$, and $I(d_{max}) \approx 0$ where d_{max} is the largest value of d that is evaluated. We advocate the use of the GLSL and CG library function `exp2`. Specifically, we use the intensity function

$$I = \exp_2(-2d^2) \quad , \tag{2}$$

which is never identically 0, but has dropped to 0.011 at $d = 1.8$. We generally find it safe to ignore values for $d > 1.8$ and set $d_{max} = 1.8$. This combination of I and d_{max} yields lines that appear thin while having very little visible aliasing.

Figure 4 shows a comparison of three methods, namely 16 times super-sampling (left), our method (in the middle) and prefiltering by direct convolution (right) with a unit variance Gaußian. The Gaußian filter was used by Chan et al. [Chan and Durand 2005]. Another popular filter is the cone function [Gupta and Sproull 1981; McNamara et al. 2000].

Both our method and Gaußian prefiltering produce lines that are visibly smoother than even 16 times super-sampling. The main difference between our method and direct convolution is that using the

---

[1] In fact, that is true of all prefiltering approaches to line drawing due to the mentioned issues related to line segments and the need to handle multiple lines that are intersecting or close together.

Figure 4: A set of lines of unit width drawn using 4×4 super sampling (left), distance filtering using $I(d) = 2^{-2d^2}$ (middle), and direct convolution with a Gaußian (right).

latter method the intersections of several lines appear to bulge a bit more. Conversely, our method produces lines that seem to contract slightly where several lines meet (cf. Figure 3). However, both phenomena are perceptible only if $I(d)$ is significantly greater than zero for very large d as in Figure 3.

3 The ID Buffer Method

The strength of the single pass method is its simplicity and the efficiency which is due to the fact that the method is single pass and requires very little work per fragment. In a generalization to non-convex polygons, more work is required per fragment since the closest polygon feature is sometimes a vertex and not a line. Unfortunately, distance to vertex is not an affine function and it must be computed in a fragment program. Moreover, N-gons for very large N are only possible to handle if they are triangulated. Thus, our generalization of the single pass method works on triangulated N-gons, and unfortunately this means that the quality depends, to some extent, on the triangulation.

Due to these issues, we propose the following method for completely general polygons. The method subsumes a technique for prefiltered line drawing which resembles the one proposed by Chan et al. [Chan and Durand 2005] but differs in important respects:

First of all, Chan et al. use the OpenGL line primitive to generate the fragments covered by the filtered line. In this work, we use a rectangle represented by a quad primitive since there are restrictions on the maximum line thickness supported by OpenGL and even more severe restrictions in the Direct3D API. Second, we compute the intensity based on the distance to the actual line *segment* whereas Chan et al. and McNamara et al. [McNamara et al. 2000] compute the intensity based on the four distances to the outer edges of the rectangle that represents the line. Our method allows for more precise rounding of the ends of the line segment as shown in Figure 6.

The position of each quad vertex is computed in a vertex program using the following formulas

$$
\begin{aligned}
\mathbf{p}_0 &= \mathbf{v}_0, & \mathbf{p}_0.xy &\mathrel{+}= \mathbf{S}\,(-\mathbf{a}-\mathbf{b})\,t\,\mathbf{v}_0.w \\
\mathbf{p}_1 &= \mathbf{v}_1, & \mathbf{p}_1.xy &\mathrel{+}= \mathbf{S}\,(+\mathbf{a}-\mathbf{b})\,t\,\mathbf{v}_1.w \\
\mathbf{p}_2 &= \mathbf{v}_1, & \mathbf{p}_2.xy &\mathrel{+}= \mathbf{S}\,(+\mathbf{a}+\mathbf{b})\,t\,\mathbf{v}_1.w \\
\mathbf{p}_3 &= \mathbf{v}_0, & \mathbf{p}_3.xy &\mathrel{+}= \mathbf{S}\,(-\mathbf{a}+\mathbf{b})\,t\,\mathbf{v}_0.w \;,
\end{aligned}
\tag{3}
$$

where $\mathbf{v}_0$ and $\mathbf{v}_1$ are the line vertices in clip coordinates. t is the line thickness in window coordinates, $\mathbf{a}$ and $\mathbf{b}$ are unit vectors along and perpendicular to the line in window coordinates and $\mathbf{S}$ is a matrix transforming from window to normalized device coordinates from

whence scaling by w brings us back to clip coordinates. The terms are illustrated in Figure 6.

The window space positions of $\mathbf{v}_i$ are passed to the fragment program where the distance from the fragment center to the line segment is computed. In fact, we only need the square distance for the intensity function (2) which is not costly to compute.


Figure 6: This figure illustrates how a single prefiltered line is drawn. The two endpoints are projected and then the line is made longer and wider in 2D.

When lines are drawn in a separate pass, the problem is to ensure that a line fragment is drawn if the line is the edge of the polygon which is drawn to the same pixel position. For this purpose, an ID buffer - which is simply a window size texture containing unique polygon IDs - is used [Markosian 2000]. Markosian uses a slightly more involved method than we do. In our implementation, we only need to sample the ID buffer in pixel centers which renders the method far simpler than if one needs to sample arbitrary locations in window space. See [Markosian 2000] for details.

The ID buffer is sufficient except in the case where the edge belongs to the silhouette of the object. If a line fragment belongs to a line along the silhouette, the pixel in the ID buffer either contains 0 (background) or the ID of a polygon below the silhouette line. In either case, the line fragment is closer than the polygon below. Hence, we can use the depth buffer without the need for offsetting to distinguish this case.

The method is described succinctly below:

Pass 1: With depth testing and depth buffer updates enabled, each polygon is rendered (filled) with a unique ID as its color attribute. The produced image is the ID buffer.

Pass 2: For each polygon, all edges are drawn using the method discussed above. A given fragment belonging to a line will be written to the framebuffer either if its ID matches the corresponding pixel in the ID buffer or if the depth test is passed. The last condition ensures that silhouette fragments are drawn. The `max` blending operator is used to ensure that if several fragments pass for a given pixel, the maximum intensity passes. Depth values are not changed during this pass.

Pass 3: At this point, the framebuffer contains alpha values and colors for the lines and depth values for the filled polygons. Now, we draw the filled, shaded polygons with the depth test set to *equal*. The polygon fragments are combined with the values in the frame buffer using blending. A window filling quad is drawn to perform blending for silhouette pixels.

While the overall approach resembles that of [Herrell et al. 1995], there are important differences: Their method was created for a

Figure 5: Left: The Happy Buddha model drawn using the single pass method. Center: The single pass method used on a quad mesh. Right: ID buffer method used to draw a gear. Hardware antialiasing (16 times) was used only on the center image (since alpha testing would otherwise introduce aliasing).

fixed function pipeline, it did not use prefiltered lines, and the loops have been inverted leading to much fewer state changes and, consequently, greater efficiency. Herrell et al. had to perform all operations for each polygon before moving on.

4 Implementations and Results

The single pass method has been implemented in CG, GLSL and HLSL. The first two implementations are based on the OpenGL API while the last implementation is based on Direct3D 10. The ID buffer method has been implemented in GLSL/OpenGL and HLSL/Direct3D 9.

We have made a number of performance tests based on the OpenGL/GLSL implementations. For all methods, we strived to use the fastest paths possible. However, in the pre-Geforce 8 series hardware that we tested, the single pass method did not allow for the use of indexed primitives for reasons we have discussed. In all other cases, we used indexed primitives. We also hinted that geometry should be stored directly in the graphics card memory either as a display list or vertex buffer object. All hardware antialiasing was disabled.

Three well known models were used, namely Newell's Teapot (converted to triangles), the Stanford Bunny, and the Stanford Happy Buddha. The models were arranged to approximately fit the viewport (of dimensions 720×576) and rotated 360 degrees around a vertical axis in 3142 steps. For each test, an average frame rate was computed based on the wall-clock time spent inside the rendering function.

A variety of hardware was used for testing. We tried to cover a wide range of graphics cards. Restricting the test to a particular machine would have been difficult for compatibility reasons. Instead, the test program was ported to the various platforms: Windows XP, Linux, and Mac OSX. The ID buffer method requires the drawing of a lot of geometry. On some systems that was too much for either a single vertex buffer or display list. Rather than using a special code path for this model, we omitted the test. Note that this problem has nothing to do with the method *per se* but only with buffer sizes and the amount of geometry.

4.1 Discussion

When judging the results, which are shown in Table 1, one should bear in mind that in order to make a fair performance comparison, no *hardware* antialiasing was enabled for any of these tests. Arguably, this lends an advantage to the offset method since hardware antialiasing is required to get an acceptable result for this type of method. Despite this, the single pass method outperforms the offset method by a factor of up to seven on the Geforce 8800 GTX. In general, the advantage of the single pass method over the offset method is greater the more recent the graphics card, and this is undoubtedly due to increasing fragment processing power. The offset method is fastest only on the Quadro FX3000 which is perhaps not surprising since this graphics card is intended for CAD users. However, the relative difference is very small for the largest mesh. It is also interesting to note that the ID buffer method, albeit the slowest, performs relatively well on the 7800GTX and 8800 GTX.

The use of geometry shaders for the single pass method makes little difference on the 8800 GTX, but it makes a great difference on the less powerful 8600M GT. The latter system is a MacBook Pro whose drivers supported the disabling of perspective correct interpolation. The 8800 GTX is a WinXP system which did not support this (at the time of testing). However, the 8800 GTX has more stream processors, and stream processing might have been a bottleneck only on the 8600M GT which entails that only this card really would benefit from the reduced load.

The single pass method is clearly much faster than the ID buffer method, and while it requires the projected polygons to be convex, it often degrades gracefully in the case of non-convex polygons as shown in Figure 7f. Nevertheless, the ID–buffer is more general since the line rendering is completely decoupled from the polygon rendering. In addition, since the single pass method draws lines as a part of polygon rendering, it is not possible to antialias the exterior side of silhouettes. Figure 7a,b show the corner of a single polygon drawn using both the single pass method and the ID buffer method. The two methods produce results which are, in fact, pixel identical on the inside, but only the ID buffer method draws the exterior part. As this issue affects only silhouette edges, it is rarely noticeable. However, for meshes consisting of few large polygons, the ID buffer method is preferable. Finally, it is possible to first draw the mesh using the single pass method and then overdraw the silhouette

	GeForce FX5900XT	Quadro FX3000	GeForce 6600GT	GeForce 7800 GTX	GeForce 8800 GTX	Geforce 8600M GT
Teapot, 6320 triangles						
offset	209.31	734.76	363.04	482.68	1548.41	439.76
single	254.22	195.81	589.59	1512.09	6226.87 (5822.38)	327.24 (808.72)
ID	22.51	23.32	87.65	263.37	550.40	86.42
Bunny, 69451 triangles						
offset	44.26	138.73	65.12	66.44	176.73	124.39
single	97.68	103.65	146.7	348.22	1506.22 (1344.25)	42.30 (380.16)
ID	4.89	5.13	19.8	61.45	114.58	15
Buddha, 1087716 triangles						
offset	3.3	9.96	4.23	4.55	11.28	11.64
single	6.56	7.07	5.31	24.19	102.74 (34.12)	2.98 (42.51)

Table 1: Performance measured in frames per second. The best performing method for each combination of mesh and graphics card is highlighted in red. For the Geforce 8800 GTX and 8600M GT the numbers in parentheses indicate the performance when geometry shaders are used.


Figure 7: Top: The difference between (a) no disabling of perspective correction, (b) perspective correction disabled using the method described in Section 2, and (c) perspective correction disabled using the Direct3D 10 API. Bottom: The corner of a polygon drawn using (d) the single pass method and (e) the ID buffer method. In (f) non-convex quads are drawn using the single pass method. Notice the fairly graceful degradation.

edges using the ID buffer method along the lines of [Sander et al. 2001].

Both proposed methods draw lines of even thickness regardless of the distance or orientation of the polygons. However, an important feature of both our methods is the fact that they are highly configurable. The line intensity is simply a function of distance, and it is possible to modulate the distance or choose a different intensity function. In Figure 8 several different line styles are shown in conjunction with the single pass method.


Figure 8: Wireframe drawing using the single pass method and a line thickness which varies as a function of the position along the edge.

Another useful effect (which is not directly possible using the offset method) is to scale the line width as a function of depth. This adds a strong depth cue as seen in Figure 5 center and right. For very dense meshes, it is also useful to fade out the wireframe as a function of distance since the shading of the filled polygons becomes much more visible in this case (see Figure 5 left). Many further possibilities present themselves. For instance, in the center image of Figure 5 alpha is also a function of distance and using alpha testing, most of the interior parts of the quadrilaterals are removed.

5 Conclusions and Future Work

In this paper we have advocated the use of prefiltered lines for wireframe rendering and demonstrated that the prefiltered lines can easily be drawn on the polygons in the pass also employed for rendering the filled polygons. This leads to a simple, single pass method which

- is far more efficent than the offset based methods,
- does not suffer from the same artifacts,
- produces lines which do not need hardware antialiasing, and
- can easily be adapted to various line styles.

Our results indicate that the gap between polygon and line rendering is widening and, in terms of performance, the strength of the single pass method is that no line primitive is required.

The ID buffer method is less efficient but handles all cases – including silhouettes – and still at interactive frame rates for meshes such as the Stanford Bunny. A trivial but useful extension of this work would be to use the ID-buffer method only for the silhouette cases. This is likely to produce flawless results in almost all conceivable cases at the expense of some complexity in the implementation.

References

AKENINE-MÖLLER, T., AND HAINES, E. 2002. *Real-Time Rendering*, 2nd ed. A.K. Peters.

BÆRENTZEN, J. A., NIELSEN, S. L., GJØL, M., LARSEN, B. D., AND CHRISTENSEN, N. J., 2006. Single-pass wireframe rendering, July. Siggraph sketches are single page refereed documents.

CHAN, E., AND DURAND, F. 2005. Fast prefiltered lines. In *GPU Gems 2*, M. Pharr, Ed. Addison Wesley, Mar.

GATEAU, S. 2007. Solid wireframe. NVIDIA Whitepaper WP-03014-001_v01, February.

GUPTA, S., AND SPROULL, R. F. 1981. Filtering edges for gray-scale displays. In *SIGGRAPH '81: Proceedings of the 8th annual conference on Computer graphics and interactive techniques*, ACM Press, New York, NY, USA, 1–5.

HERRELL, R., BALDWIN, J., AND WILCOX, C. 1995. High-quality polygon edging. *IEEE Comput. Graph. Appl. 15*, 4, 68–74.

ISENBERG, T., FREUDENBERG, B., HALPER, N., SCHLECHTWEG, S., AND STROTHOTTE, T. 2003. A developer's guide to silhouette algorithms for polygonal models. *IEEE Computer Graphics and Applications 23*, 4, 28–37.

MARKOSIAN, L. 2000. *Art-based Modeling and Rendering for Computer Graphics*. PhD thesis, Brown University.

MCNAMARA, R., MCCORMACK, J., AND JOUPPI, N. P. 2000. Prefiltered antialiased lines using half-plane distance functions. In *HWWS '00: Proceedings of the ACM SIGGRAPH/EUROGRAPHICS workshop on Graphics hardware*, ACM Press, New York, NY, USA, 77–85.

MCREYNOLDS, T., AND BLYTHE, D., 1999. Advanced graphics programming techniques using opengl. SIGGRAPH '99 Course.

ROSE, D., AND ERTL, T. 2003. Interactive visualization of large finite element models. In *Proceedings of Vision Modeling and Visualization 2003*, Akademische Verlagegesellschaft, T. Ertl, B. Girod, G. Greiner, H. Niemann, H.-P. Seidel, E. Steinbach, and R. Westermann, Eds., 585–592.

SANDER, P. V., HOPPE, H., SNYDER, J., AND GORTLER, S. J. 2001. Discontinuity edge overdraw. In *SI3D '01: Proceedings of the 2001 symposium on Interactive 3D graphics*, ACM Press, New York, NY, USA, 167–174.

SEGAL, M., AND AKELEY, K. 2004. *The OpenGL Graphics System: A Specification (Version 2.0)*. SGI.

WANG, W., CHEN, Y., AND WU, E. 1999. A new method for polygon edging on shaded surfaces. *J. Graph. Tools 4*, 1, 1–10.

Appendix: Shaders

In the following, we list the vertex, geometry, and fragment shaders for the variation of the single pass method which employs the geometry shader. Note that almost all complexity is in the geometry shader.

```
// ----------------- Vertex Shader -------------------------------
#version 120
#extension GL_EXT_gpu_shader4 : enable
void main(void)
{
   gl_Position =  ftransform();
}

// ----------------- Geometry Shader -------------------------------
#version 120
#extension GL_EXT_gpu_shader4 : enable

uniform vec2 WIN_SCALE;
noperspective varying vec3 dist;
void main(void)
{
  vec2 p0 = WIN_SCALE * gl_PositionIn[0].xy/gl_PositionIn[0].w;
  vec2 p1 = WIN_SCALE * gl_PositionIn[1].xy/gl_PositionIn[1].w;
  vec2 p2 = WIN_SCALE * gl_PositionIn[2].xy/gl_PositionIn[2].w;

  vec2 v0 = p2-p1;
  vec2 v1 = p2-p0;
  vec2 v2 = p1-p0;
  float area = abs(v1.x*v2.y - v1.y * v2.x);

  dist = vec3(area/length(v0),0,0);
  gl_Position = gl_PositionIn[0];
  EmitVertex();

  dist = vec3(0,area/length(v1),0);
  gl_Position = gl_PositionIn[1];
  EmitVertex();

  dist = vec3(0,0,area/length(v2));
  gl_Position = gl_PositionIn[2];
  EmitVertex();

  EndPrimitive();
}

// ----------------- Fragment Shader -------------------------------
#version 120
#extension GL_EXT_gpu_shader4 : enable

noperspective varying vec3 dist;
const vec4 WIRE_COL = vec4(1.0,0.0,0.0,1);
const vec4 FILL_COL = vec4(1,1,1,1);

void main(void)
{
float d = min(dist[0],min(dist[1],dist[2]));
  float I = exp2(-2*d*d);
  gl_FragColor = I*WIRE_COL + (1.0 - I)*FILL_COL;
}
```